How much of a 3D scene should a language model read?

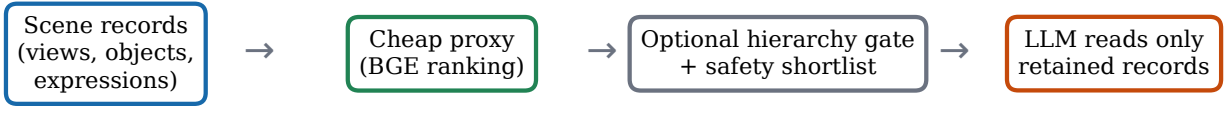

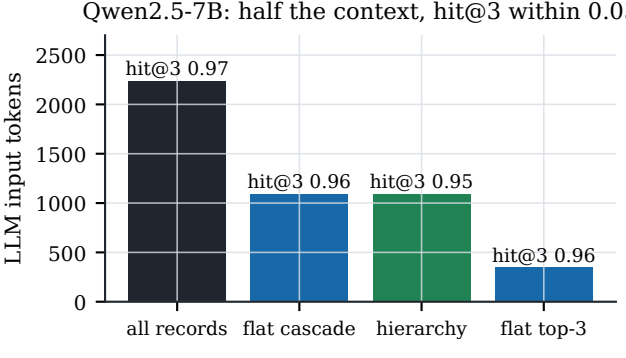

Findings

- Answer accuracy 0.61 -> 0.73 with 3 records
- Hierarchy = flat cascade at matched budget
- Safety shortlist makes the hierarchy interchangeable
- At 155k records, 96.6% of flat misses are other-scene
- Replica GT orientation error found and corrected